

Methods and Tools

Interactive Lectures

All lectures in the course will be interactive

They contain running code, as well as theory!

- Presented and discussed in frontal lectures...
- ...You can download PDFs
- ...But you will also be able to make changes and experiment

From a software perspective, the workhorses of this approach are:

- Jupyter notebooks for the presentation & interaction
- Docker containers for the setup and distribution


Both are widely used systems:

- Jupyter is a user favorite when it comes to data science
- Docker is a state-of-the-art system for managing services

A Few Words about Docker

Docker

Docker is a system for running software in "containers"




Think of a container as a **lightweight virtual machine**:

- (Essentially) the same level of isolation
- ...But smaller disk footprint, faster setup and operation, etc.

Docker

Docker is a system for running software in "containers"



Using containers has many advantages:

- Multiple environments on the same machine
- Improved isolation, robustness, and reproducibility
- Easier replication (scalability of cloud services)
- ...

Docker

During this course we will see many problems

...And tackle them with **many techniques**:

- Classical Machine Learning
- Deep Learning
- Statistics
- Signal processing
- Declarative optimization
- Differential Equations
- Agent based simulation
- ...

Managing dependences would quickly become hellish

Docker

With docker, we can simply use a different container per case study

Inside each container we will have:

- All the needed libraries & tools
- A running instance of a Jupyter server

In the host machine (your PC):

- We will just open a browser...
- ...And connect to the Jupyter server

Two key concepts in Docker

- A container is a (sort of) running, lightweight, Virtual Machine
- An image is (sort of) the content of the hard disk of the VM

The image can be used to instantiate multiple containers

Building an Image

Images in docker are built by:

- Starting from a base image on [Docker Hub](#)
- Copying content between the host and the container
- Running commands in the container

The process is controlled via a Dockerfile

- Just a text file with a specific syntax
- There is an [extensive reference](#), but we only care about a few commands

To build an image, we can use:

```
docker build .
```

...From the directory with the Dockerfile

An Example Dockerfile

This is a **minimal Dockerfile** for this lecture:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The `FROM` keyword specifies the base image

An Example Dockerfile

This is a **minimal Dockerfile** for this lecture:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The `RUN` keyword runs a command
- In our case, we install a number of python packages

An Example Dockerfile

This is a **minimal Dockerfile** for this lecture:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The `COPY` keyword transfers data from the host to the container
- The first path refers to the host
- The second path to the container

An Example Dockerfile

This is a **minimal Dockerfile** for this lecture:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The `WORKDIR` changes the current directory in the container
- It's like running `cd` in the container

An Example Dockerfile

This is a **minimal Dockerfile** for this lecture:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The `CMD` keyword is triggered only when we **run** a container
- It's the first command that the container should execute
- It does nothing when building an image

An Example Dockerfile

When we run `docker build .` for our file:

- The docker daemon downloads the base image, if not already available
- A container is started
- All operations in the Dockerfile are executed
- The resulting container is dumped, to create an **image**

You can check that a new image has been built using:

```
docker image ls
```

You will see an entry with no name:

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	96b910c1514f	3 seconds ago	1.36GB

An Example Dockerfile

You can assign a name to an image using:

```
docker build . -t <name of the image>
```

You can remove an image with:

```
docker image rm <image name or id>
```

- Useful to free space, however...
- ...Images are incremental! Docker stores only the differences
- ...So, don't worry too much about space usage

You can remove all images with no running container with:

```
docker image prune
```

Running a Container

You can instantiate and run a container with:

```
docker run <image name or id>
```

- The container `stdout` will be piped (i.e. connected) to your terminal
- By default, this is not the case for `stdin`
- You can make the container interactive with the `-it` options
- You can autoremove the container at the end with `--rm`
- You can sync folders in the host and on the container using volumes

The documentation is extensive

Running a Container

You can obtain the list of all containers with:

```
docker ps
```

- The option `-a` shows all containers (incl. those that are stopped)

You can remove a container with

```
docker rm <container id>
```

- As you see, it's a very flexible system
- ...But also a bit complex

That's why we will automate most operations using Docker Compose

Docker Compose

Docker Compose is a tool to help the management of containers

In a second `docker-compose.yml` file, you specify:

- Which "services" (i.e. container) should be built and run
- How to build them
- Which options to use when running them
- ...

All in a human-readable, declarative format

A Docker Compose Example

Let's see the `docker-compose.yml` for this lecture:

```
version: '2.0'
services:
  jupyter:
    build: .
    ports:
      - "8888:8888"
    volumes:
      - ./app
```

- `version` refers to the Docker Compose syntax
- `services` is followed by a list of the containers
- `jupyter` is our service
- `build` specifies where the `Dockerfile` can be found

A Docker Compose Example

Let's see the `docker-compose.yml` for this lecture:

```
version: '2.0'
services:
  jupyter:
    build: .
    ports:
      - "8888:8888"
    volumes:
      - .:/app
```

- `ports` tells which ports to expose to `docker run`
- `volumes` specifies which folders to sync
- In our case `"/app/notebooks"` on the container
- ...Will actually be `"/notebooks"` on the host

Benefits of Using Docker Compose

We need to use one more tool, but now we can:

Build and run a container with:

```
docker-compose up
```

- The command can also restart a stopped container

Stop the container with CTRL+C, or with:

```
docker-compose stop
```

Stop and remove the container with:

```
docker-compose down
```

...Which is considerably simpler than before!

Our Jupyter Setup

Our Setup

We will often work with this development setup

The folder with the notebooks is structured as follows:

```
notebook1.pynb
notebook2.pynb
...
util      <-- module
assets    <-- images and such
rise.css  <-- for the "slide" mode
```

Our Setup

We will often work with this development setup

The folder with the notebooks is structured as follows:

```
notebook1.pynb
notebook2.pynb
...
util +-- __init__.py
      +-- XYZ.py      <-- submodule
      +-- YZX.py      <-- submodule
      +-- ...
font
rise.css
```

Most important part: the use of **a module** besides notebooks

Our Setup

Working with modules provides some advantages:

We do not need to keep all our code in the notebooks. We can:

- Share functions **between cells**
- Share functions **between notebooks**
- IDEs can offer **more functionality** if they recognize a module

...But also a significant disadvantage:

- Python modules are compiled first when loaded...
- ...The loaded version is **not updated** when the source changes

This is very inconvenient at development time

Our Setup

We can circumvent this thanks to Jupyter "magic" extensions

The first one is the "autoreload" extension

```
In [1]: %load_ext autoreload
        %autoreload 2
```

- `load_ext` will enable the extension
- `autoreload 2` will reload all modules before code execution

This is *inefficient, but convenient* during development

- Together with the use of volumes (in docker-compose)...
- ...This allows us to update the code without re-building the docker image

Starting a Notebook

Let's look back to the `CMD` keyword in our `Dockerfile`:

```
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", \  
      "--ip=0.0.0.0", "--allow-root"]
```

This is translated to:

```
jupyter notebook --port=8888 --no-browser --ip=0.0.0.0 --allow-root
```

- `--port 8888`: the server listen on port 8888
- `--no-browser`: do not open the browser (there's no browser in the container)
- `--ip=0.0.0.0`: listen on all network interfaces
- `--allow-root`: we operate as `root` (admin) on the container

Starting a Notebook

When we run:

```
docker-compose up
```

The output will look like:

```
Starting ad_stat_jupyter_1 ... done
...
...Use Control-C to stop this server and shut down all kernels...
...
...To access the notebook...
...copy and paste one of these URLs:
...   http://34b908cf2362:8888/?token=82e337a2be9915cdebce276bf...
... or http://127.0.0.1:8888/?token=82e337a2be9915cdebce276bf...
```

- The last URL can be **copy-pasted** in your favorite browser

Starting a Notebook

When we run:

```
docker-compose up
```

The output will look like:

```
Starting ad_stat_jupyter_1 ... done
...
...Use Control-C to stop this server and shut down all kernels...
...
...To access the notebook...
...copy and paste one of these URLs:
...   http://34b908cf2362:8888/?token=82e337a2be9915cdebce276bf...
... or http://127.0.0.1:8888/?token=82e337a2be9915cdebce276bf...
```

- The `token` is cached by the browser and grants access to the notebooks

Our Setup

- We will also use the `ipymp1` package and the `widget` jupyter magic
- This will display basic tools to rescale and zoom images

```
In [5]: %matplotlib widget  
from matplotlib import pyplot as plt  
import numpy as np  
x = np.linspace(0, 2*np.pi, 100)  
plt.figure(figsize=(9, 3))  
plt.plot(x, np.sin(x))  
plt.tight_layout()
```

