

# KDE for Our Anomaly Detection Case Study

---

## A Few Tweaks

**We will work with log probabilities**

- This is what sklearn does by default
- ...And simplifies some operations
- I.e. products become sums

**Additionally, we will work with negated (log) probabilities**

- They can be interpreted as alarm signals
- ...Which is the customary approach in anomaly detection

**Overall, our anomaly detection condition becomes:**

$$-\log f(x, \omega) \geq \theta$$

...Which is equivalent to the previous formulation

# Training and Testing

## We will split our data in two segments

A **training set**:

- This will include only data about the normal behavior
- Ideally, there should be no anomalies here (we do not want to learn them!)
- We will use it to fit a KDE model

A **test set**:

- To assess how well the approach can generalize

## If the training set contains some anomalies

- Things may still be fine, as long as they are very infrequent
- ...Since we will still learn that they have low probability

# Training and Testing

**In time series data sets are often split chronologically:**

```
In [4]: train_end = pd.to_datetime('2014-10-24 00:00:00')  
  
nab.plot_series(data, labels, test_start=train_end)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

- Green: training set, orange: test set

# Normalization

## We start by normalizing the data

- We isolate the training data
- We can use a date as a separator since we have a `datetime` index

```
In [5]: tmp = data[data.index < train_end]
```

- Then we compute the maximum on the training data:

```
In [6]: trmax = tmp['value'].max()
```

- Finally, we normalize all data
- ...And separate the normalized training data

```
In [7]: data['value'] = data['value'] / trmax  
data_tr = data[data.index < train_end]
```

# Normalization

Here we see the result:

```
In [9]: nab.plot_series(data, labels, windows)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

Same as before, just with a different scale

# Fitting the Estimator

## Now, we need to estimate the bandwidth

Since we are in a univariate case, we will use the rule of thumb:

```
In [12]: q1 = data_tr['value'].quantile(0.25)
q3 = data_tr['value'].quantile(0.75)
sigma = data_tr['value'].std()
m = len(data_tr)
h = 0.9 * min(sigma, (q3-q1) / 1.34) * m**(-0.2)
print(f'The estimated bandwidth is {h:.3f}')
```

The estimated bandwidth is 0.035

## Then we fit a univariate KDE estimator

- We will use the scikit-learn library

```
In [13]: kde = KernelDensity(kernel='gaussian', bandwidth=h)
kde.fit(data_tr);
```

# Fitting the Estimator

Let's have a look at the estimate distribution:

```
In [15]: vmax = data['value'].max()
xr = np.linspace(0, vmax, 100)
nab.plot_density_estimator_1D(kde, xr)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.



# Alarm Signal

We can now obtain (and plot) our alarm signal:

```
In [17]: ldens = kde.score_samples(data) # Obtain log probabilities
         signal = pd.Series(index=data.index, data=-ldens) # Build series with neg. prob.
         nab.plot_series(signal, labels=labels, windows=windows) # Plot
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

- Noticed how the the process (inference) was noticeably slower than training?
- Can you tell why this was the case?

# Detecting Anomalies

Let us pick a threshold (at random, for now) and try to detect some anomalies:

```
In [18]: thr = 4
pred = pd.Series(signal.index[signal >= thr])
print(pred)

0    2014-07-03 19:00:00
1    2014-09-06 22:30:00
2    2014-09-06 23:00:00
3    2014-11-02 01:00:00
4    2014-11-02 01:30:00
5    2015-01-01 01:00:00
Name: timestamp, dtype: datetime64[ns]
```

- We just apply the filter `signal >= thr` to the signal index
- This yields a number of (predicted) anomalous timestamps
- `pred` is pandas `series` with the detected anomalous timestamps
- Our module contains a function for this step:

# Detecting Anomalies

Let us plot our predictions on the series:

```
In [19]: nab.plot_series(signal, labels=predictions, windows=pred)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

- Not bad, especially for a random attempt!
- There are a few false positives, which are **very common** in anomaly detection

# Metrics for Anomaly Detection

---

# Metrics for Anomaly Detection

**Evaluating the quality of an Anomaly Detection system can be tricky**

- Usually, we do not need to match the anomalies exactly
- Sometimes we wish to anticipate anomalies
- ...But sometimes we just want to detect them in past data

There is no "catch-all" metric, like accuracy in classification

**It is much better to devise a cost model**

- We evaluate the cost and benefits of our predictions:
- By doing this, we focus on the value for our customer

**This is important for all industrial problems!**

# A Simple Cost Model

## We will use a simple cost model

Remember that our goals are:

- Analyzing anomalies
- Anticipating anomalies

First, we define:

- **True Positives** as windows for which we detect at least one anomaly
- **False Positives** as detected anomalies that do not fall in any window
- **False negatives** as anomalies that go undetected
- **Advance** as the time between an anomaly and when first we detect it

# A Simple Cost Model

They can be computed using a function in our module:

```
def get_metrics(pred, labels, windows)
```

```
In [20]: tp, fp, fn, adv = nab.get_metrics(pred, labels, windows)
print('True Positives', tp)
```

```
True Positives 0    2014-11-01 19:00:00
1    2015-01-01 01:00:00
dtype: datetime64[ns]
```

```
In [21]: print('False Positives', fp)
```

```
False Positives 0    2014-07-03 19:00:00
1    2014-09-06 22:30:00
2    2014-09-06 23:00:00
dtype: datetime64[ns]
```

# A Simple Cost Model

They can be computed using a function in our module:

```
def get_metrics(pred, labels, windows)
```

```
In [22]: print('False Negatives', fn)
```

```
False Negatives 0    2014-11-27 15:30:00
1    2014-12-25 15:00:00
2    2015-01-27 00:00:00
dtype: datetime64[ns]
```

```
In [23]: print('Advance', adv)
```

```
Advance 0    -1 days +18:00:00
1         0 days 00:00:00
dtype: timedelta64[ns]
```



# A Simple Cost Model

Then we introduce:

- A cost  $c_{alarm}$  for loosing time in analyzing false positives
- A cost  $c_{missed}$  for missing an anomaly
- A cost  $c_{late}$  for a late detection (partial loss of value)

Our cost is then given by:

```
In [24]: c_alarm = 1 # Cost of investigating a false alarm
c_missed = 10 # Cost of missing an anomaly
c_late = 5 # Cost for late detection

# A simple cost model
cost = c_alarm * len(fp) + \
      c_missed * len(fn) + \
      c_late * (len([a for a in adv if a.total_seconds() <= 0]))
print(f'The cost with the current predictions is: {cost}')
```

The cost with the current predictions is: 43

# Threshold Choice

---

# Refactoring the Cost Model

We will use this cost model for **choosing our threshold**

With this aim, we encapsulate the formula in a Python class in our module:

```
class ADSimpleCostModel:
    def __init__(self, c_alarm, c_missed, c_late):
        self.c_alarm = c_alarm
        self.c_missed = c_missed
        self.c_late = c_late

    def cost(self, signal, labels, windows, thr):
        pred = get_pred(signal, thr)
        tp, fp, fn, adv = get_metrics(pred, labels, windows)
        adv_det = [a for a in adv if a.total_seconds() <= 0]
        cost = self.c_alarm * len(fp) + \
            self.c_missed * len(fn) + \
            self.c_late * (len(adv_det))
        return cost
```

# Refactoring the Cost Model

**We can now evaluate our cost by:**

- Instantiating the class with our values for  $C_{alarm}$ ,  $C_{missed}$ ,  $C_{late}$
- Calling the `cost` function with different threshold values

```
In [25]: cmodel = nab.ADSimpleCostModel(c_alarm, c_missed, c_late)
cost = cmodel.cost(signal, labels, windows, thr)
print(f'The cost with the current predictions is: {cost}')
```

```
The cost with the current predictions is: 43
```

**The process is rather efficient:**

- When we change the threshold
- ...We do not need to rebuild the alarm signal

So, no need to repeat either training or inference

# Effect of the Threshold

**We can now look (over all data) at the effect of changing the threshold:**

- This is the **response surface** or **cost function landscape**
- ...For an **idealized problem** (i.e. over all the data)

```
In [27]: cmodel = nab.ADSimpleCostModel(c_alarm, c_missed, c_late)
thr_range = np.linspace(1, 10, 100)
cost_range = [cmodel.cost(signal, labels, windows, thr) for thr in thr_range]
cost_range = pd.Series(index=thr_range, data=cost_range)
nab.plot_series(cost_range)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

# Choosing the Threshold

**Ideally, we wish to choose the best threshold**

However, we have two problems:

- We cannot really use all the data
- ...we need to keep some away to test generalization
- The threshold-to-cost function is non-smooth and non-differentiable

**Luckily, both are easy to address**

We can:

- Define a validation set to use when optimizing the threshold
- Use a simple line search approach, since  $\theta$  is a scalar

# Define a Validation Set

**First, we define our validation set:**

- It should contain some anomalies
- ...But some should be left out (for testing)

```
In [28]: val_end = pd.to_datetime('2014-12-10 00:00:00')
nab.plot_series(data, labels, val_start=train_end, test_start=val_end)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

# Optimize the Threshold

**We will use a line sampling approach:**

We build the validation set and define a range of "sampled" thresholds:

```
In [29]: signal_opt = signal[signal.index < val_end]
labels_opt = labels[labels < val_end]
windows_opt = windows[windows['end'] < val_end]
thr_range = np.linspace(0, 10, 100)
```

- Notice that we included the training data in the validation set
- ...This is ok, since we are not using the validation set to prevent overfitting

**Then, we use a function from our module to pick the best one:**

```
def opt_thr(signal, labels, windows, cmodel, thr_range):
    costs = [cmodel.cost(signal, labels, windows, thr)
              for thr in thr_range]
    costs = np.array(costs)
    best_idx = np.argmin(costs)
    return thr_range[best_idx], costs[best_idx]
```



# Optimize the Threshold

Let's see the results:

```
In [30]: best_thr, best_cost = nab.opt_thr(signal_opt, labels_opt,
                                           windows_opt, cmodel, thr_range)
print(f'Best threshold: {best_thr}, corresponding cost: {best_cost}')
```

Best threshold: 4.747474747474747, corresponding cost: 15

- The reported cost is on the training set

**For all the data (yes, we are cheating a bit) we have:**

```
In [23]: ctst = cmodel.cost(signal, labels, windows, best_thr)
print(f'Cost on the whole dataset {ctst}')
```

Cost on the whole dataset 45

- This is, as expected, **suboptimal**
- ...But it works! We have our **first complete anomaly detection approach**